

Batch: B - 1 Roll No.: 16014022050

Experiment No. 6

Title : Implementation and Analysis of Perceptron Learning Algorithm

Course Outcome:

CO3: Differentiate between types of machine learning and implement supervised and unsupervised learning

Books/ Journals/ Websites referred:

1. “Artificial Intelligence: a Modern Approach” by Russell and Norving, Pearson education Publications
2. “Artificial Intelligence” By Rich and knight, Tata McGraw Hill Publications

Introduction:

The perceptron is one of the earliest and simplest supervised learning algorithms for binary classification. It attempts to find a linear decision boundary that separates data points of different classes. The algorithm adjusts its weights iteratively based on misclassified examples until convergence or until a fixed number of epochs.

In this experiment, we use the **Iris dataset** (a popular dataset with 150 flower samples from three species: Setosa, Versicolor, and Virginica). Since the perceptron works only for binary classification, we restrict our analysis to **two classes (Setosa vs. Versicolor)**. We will analyze the effect of varying the number of epochs on model accuracy and visualize the learning process with multiple graphs.

Implementation:

Part 1: Install & Import Libraries

```
[1] # Make sure libraries are available
!pip install scikit-learn matplotlib numpy

# Imports
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

→ Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (1.6.1)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.16.2)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.60.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyParsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.4)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib) (1.17.0)
```

Part 2: Load the Iris Dataset

```
[2] # Load Iris dataset from sklearn
iris = load_iris()

# Use only first 100 samples (Setosa & Versicolor) → binary classification
X = iris.data[:100, :2] # only first 2 features for easy visualization
y = iris.target[:100]

# Convert labels {0,1} → {-1,1}
y = np.where(y == 0, -1, 1)

print("Shape of X:", X.shape)
print("Shape of y:", y.shape)
print("Classes in y:", np.unique(y))

→ Shape of X: (100, 2)
Shape of y: (100,)
Classes in y: [-1  1]
```

Part 3: Train-Test Split & Standardization

```
[3] # Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Standardize features (better convergence)
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)

print("Training samples:", X_train.shape[0])
print("Testing samples:", X_test.shape[0])

→ Training samples: 80
Testing samples: 20
```

Part 4: Define Perceptron Class

```
[4]  ✓ 0s
    class Perceptron:
        def __init__(self, learning_rate=0.01, n_epochs=10):
            self.lr = learning_rate
            self.n_epochs = n_epochs
            self.weights = None
            self.bias = 0
            self.errors_ = []

        def fit(self, X, y):
            self.weights = np.zeros(X.shape[1])
            for _ in range(self.n_epochs):
                errors = 0
                for xi, target in zip(X, y):
                    update = self.lr * (target - self.predict(xi))
                    self.weights += update * xi
                    self.bias += update
                    errors += int(update != 0.0)
                self.errors_.append(errors)
            return self

        def net_input(self, X):
            return np.dot(X, self.weights) + self.bias

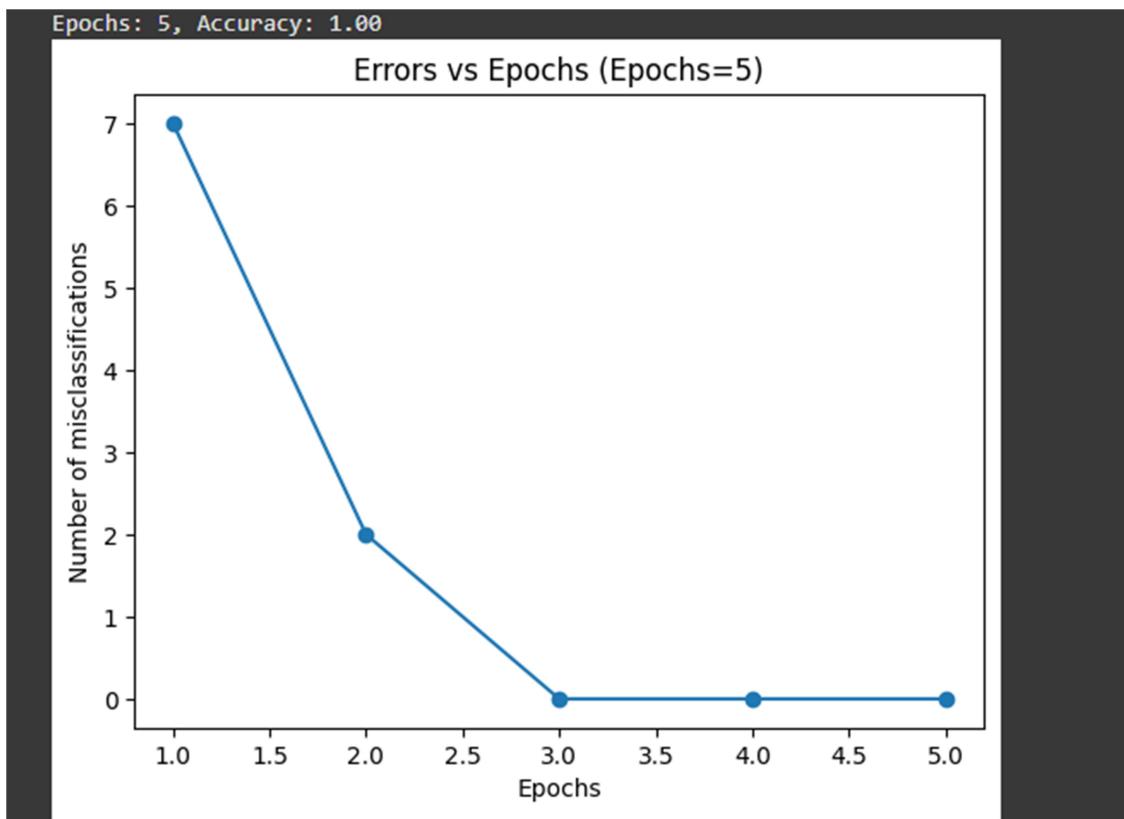
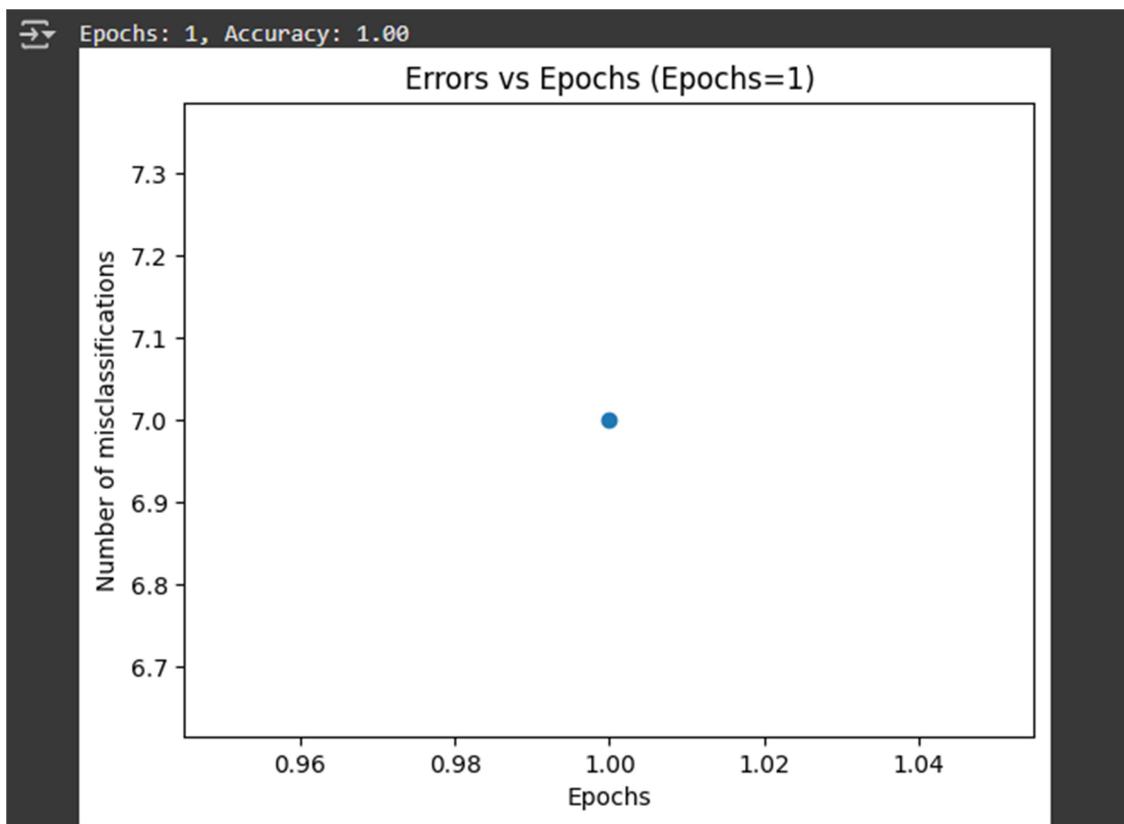
        def predict(self, X):
            return np.where(self.net_input(X) >= 0.0, 1, -1)
```

Part 5: Train with Different Epochs & Plot

```
[5]  ✓ 2s
    epoch_list = [1, 5, 10, 20, 50]
    accuracies = []

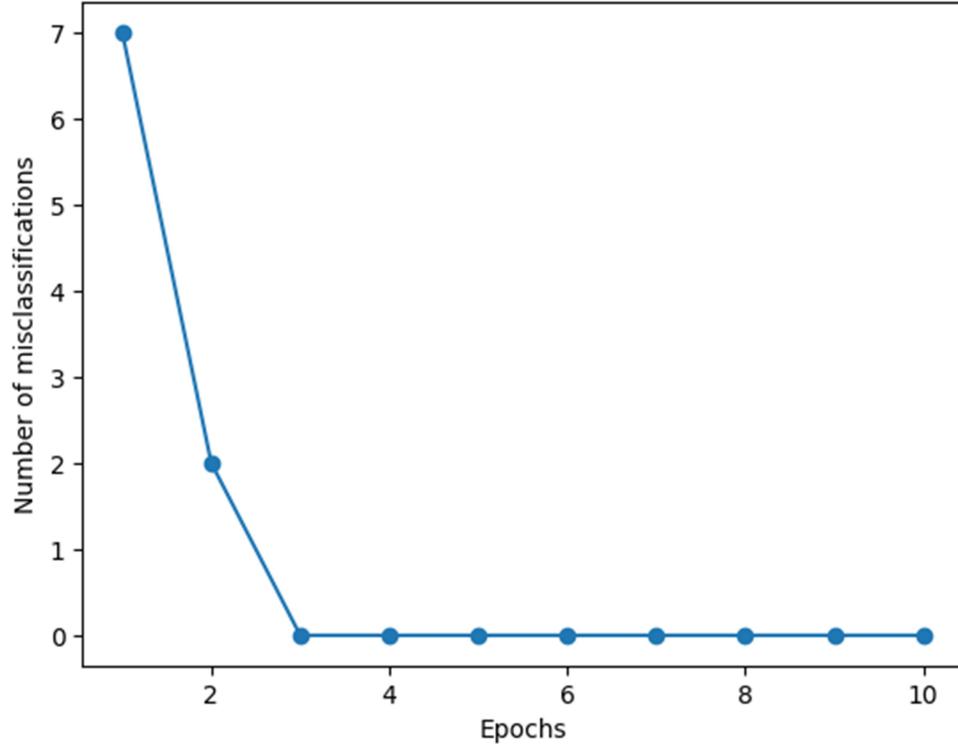
    for ep in epoch_list:
        ppn = Perceptron(learning_rate=0.01, n_epochs=ep)
        ppn.fit(X_train, y_train)
        y_pred = ppn.predict(X_test)
        acc = np.mean(y_pred == y_test)
        accuracies.append(acc)
        print(f"Epochs: {ep}, Accuracy: {acc:.2f}")

        # Misclassification plot
        plt.plot(range(1, ep+1), ppn.errors_, marker='o')
        plt.title(f"Errors vs Epochs (Epochs={ep})")
        plt.xlabel("Epochs")
        plt.ylabel("Number of misclassifications")
        plt.show()
```



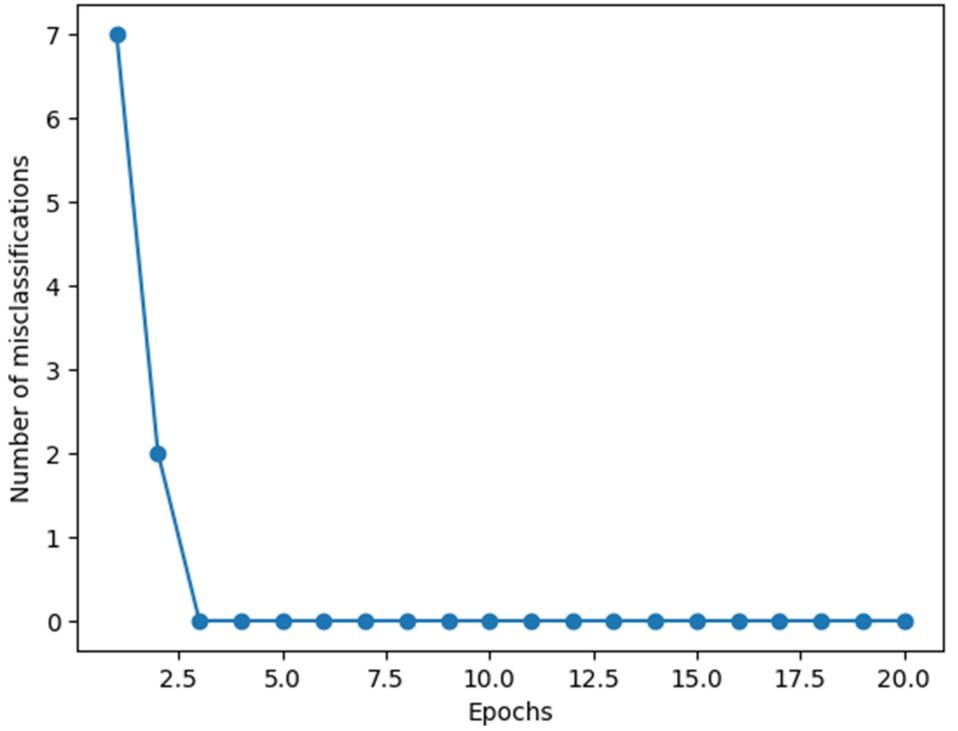
Epochs: 10, Accuracy: 1.00

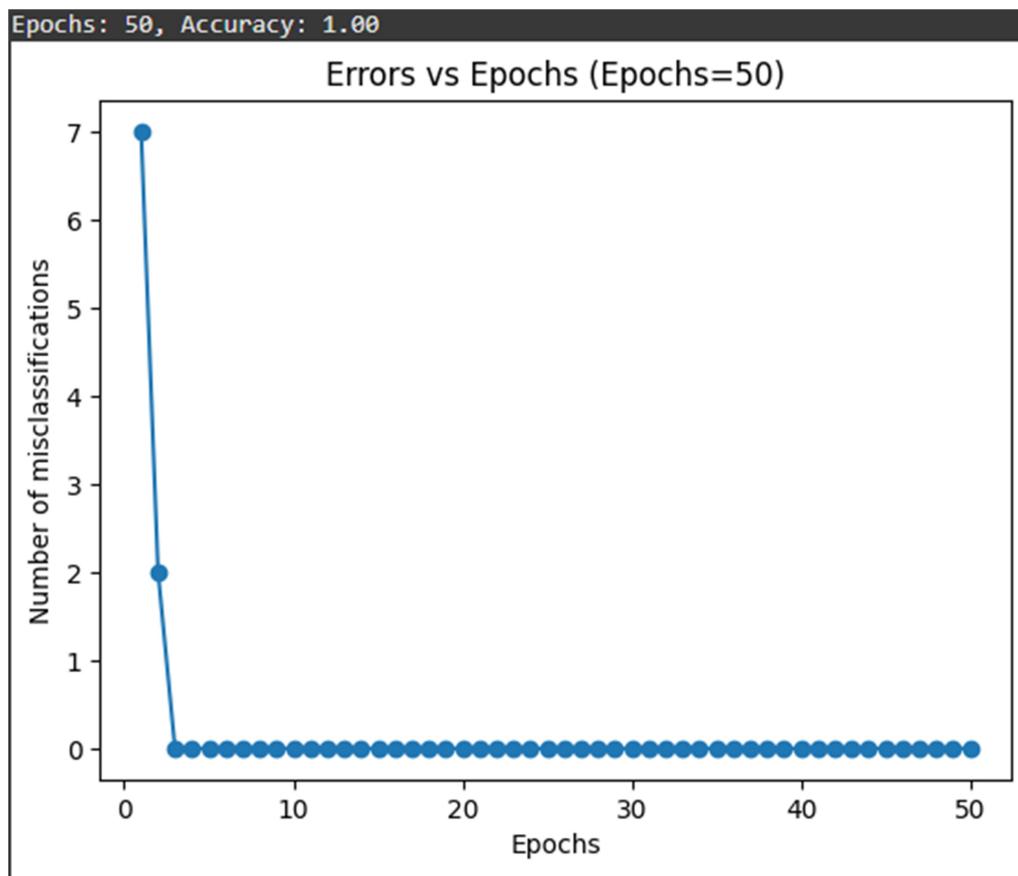
Errors vs Epochs (Epochs=10)



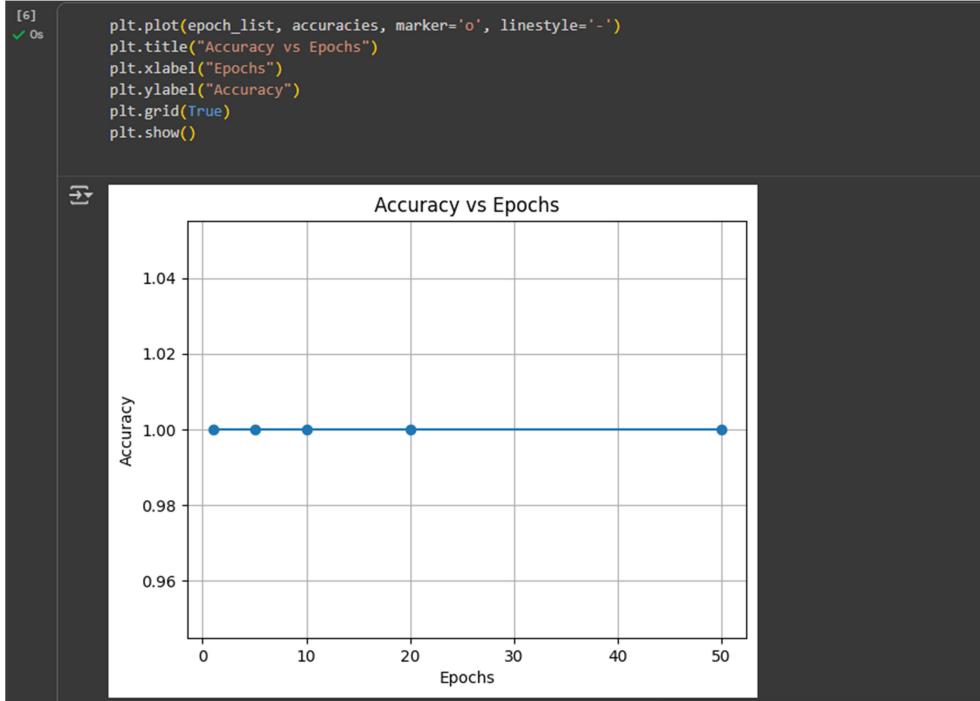
Epochs: 20, Accuracy: 1.00

Errors vs Epochs (Epochs=20)





Part 6: Accuracy vs Epochs Graph



Part 7: Decision Boundary Visualization

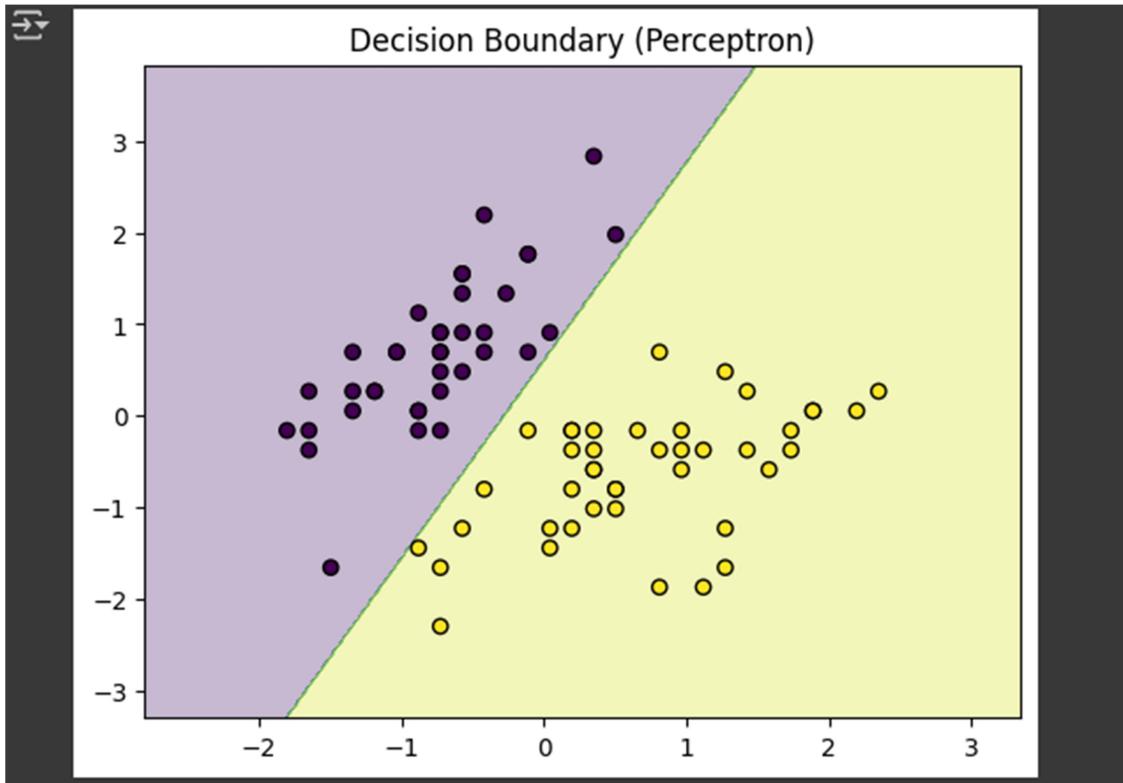
```
[7]
✓ 0s

# Train a final perceptron with 20 epochs
ppn = Perceptron(learning_rate=0.01, n_epochs=20)
ppn.fit(X_train, y_train)

# Create grid
x_min, x_max = X_train[:, 0].min() - 1, X_train[:, 0].max() + 1
y_min, y_max = X_train[:, 1].min() - 1, X_train[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                      np.arange(y_min, y_max, 0.01))

Z = ppn.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

# Plot
plt.contourf(xx, yy, Z, alpha=0.3)
plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train, edgecolor='k')
plt.title("Decision Boundary (Perceptron)")
plt.show()
```



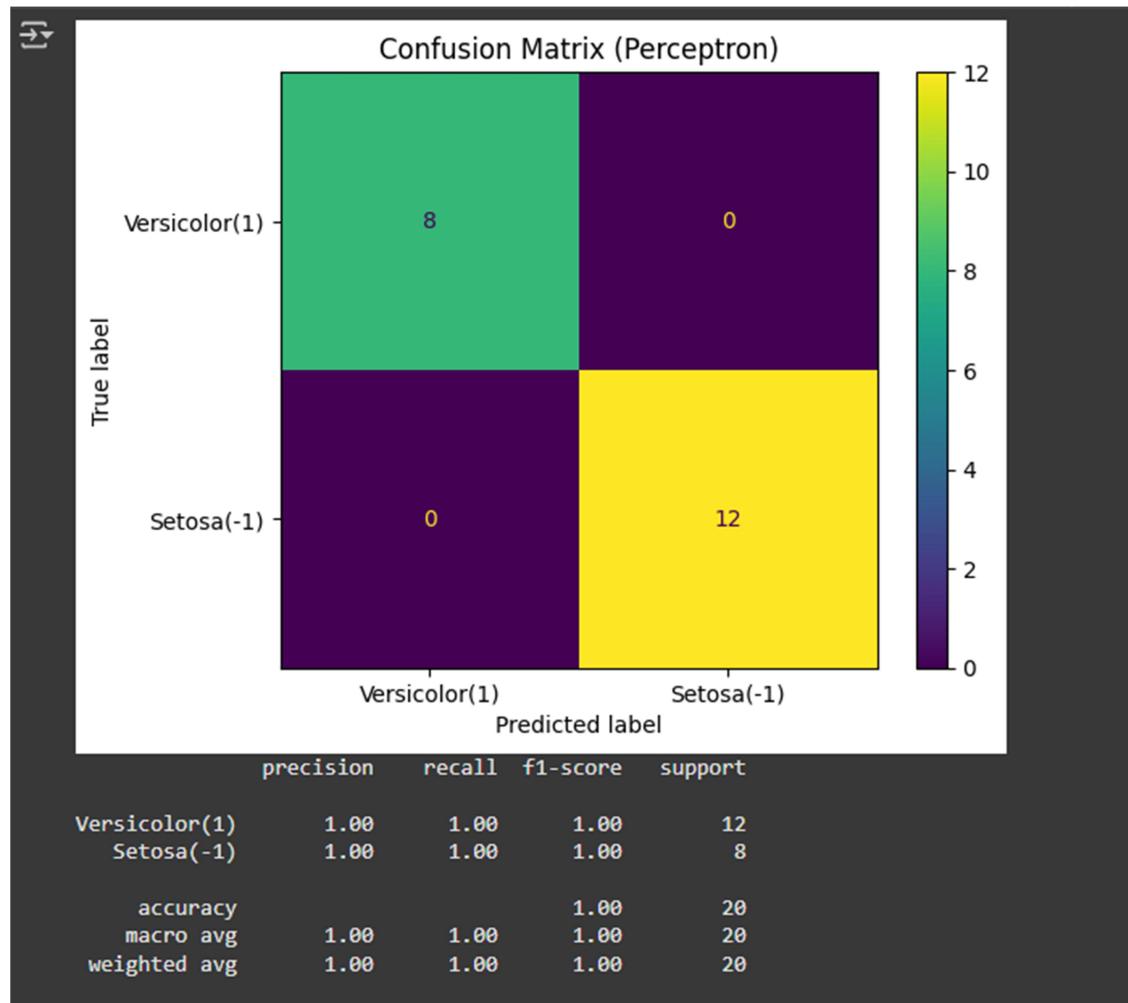
Part 8: Confusion Matrix & Classification Report

```
[8] 1s
from sklearn.metrics import confusion_matrix, classification_report, ConfusionMatrixDisplay

ppn = Perceptron(learning_rate=0.01, n_epochs=20)
ppn.fit(X_train, y_train)
y_pred = ppn.predict(X_test)

# Confusion matrix
cm = confusion_matrix(y_test, y_pred, labels=[1,-1])
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=['Versicolor(1)', 'Setosa(-1)'])
disp.plot()
plt.title('Confusion Matrix (Perceptron)')
plt.show()

# Text report
print(classification_report(y_test, y_pred, target_names=['Versicolor(1)', 'Setosa(-1)']))
```



Part 9: Printing accuracy

```
[13]: 0s
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
epoch_list = [1, 5, 10, 20, 50]

for ep in epoch_list:
    ppn = Perceptron(learning_rate=0.01, n_epochs=ep)
    ppn.fit(X_train, y_train)

    # Predictions
    y_train_pred = ppn.predict(X_train)
    y_test_pred = ppn.predict(X_test)

    # Train metrics
    train_acc = accuracy_score(y_train, y_train_pred)
    train_prec = precision_score(y_train, y_train_pred, pos_label=1)
    train_rec = recall_score(y_train, y_train_pred, pos_label=1)
    train_f1 = f1_score(y_train, y_train_pred, pos_label=1)

    # Test metrics
    test_acc = accuracy_score(y_test, y_test_pred)
    test_prec = precision_score(y_test, y_test_pred, pos_label=1)
    test_rec = recall_score(y_test, y_test_pred, pos_label=1)
    test_f1 = f1_score(y_test, y_test_pred, pos_label=1)

    # Print nicely
    print(f"\nEpochs: {ep}")
    print(f" Train -> Accuracy: {train_acc:.2f}, Precision: {train_prec:.2f}, Recall: {train_rec:.2f}, F1: {train_f1:.2f}")
    print(f" Test -> Accuracy: {test_acc:.2f}, Precision: {test_prec:.2f}, Recall: {test_rec:.2f}, F1: {test_f1:.2f}")


[13]: 0s
```

```
→ Epochs: 1
Train -> Accuracy: 0.99, Precision: 0.98, Recall: 1.00, F1: 0.99
Test -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00

Epochs: 5
Train -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00
Test -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00

Epochs: 10
Train -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00
Test -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00

Epochs: 20
Train -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00
Test -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00

Epochs: 50
Train -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00
Test -> Accuracy: 1.00, Precision: 1.00, Recall: 1.00, F1: 1.00
```

Post Lab Descriptive Questions:

1. What is the perceptron learning rule, and how does it update weights?

The perceptron learning rule updates weights based on classification errors. If a sample is misclassified, the weight vector is adjusted in the direction of the input. Update rule:

$$w_{new} = w_{old} + \eta(y - \hat{y})x$$

$$b_{new} = b_{old} + \eta(y - \hat{y})$$

where:

- η = learning rate
- y = true label
- \hat{y} = predicted label

2. Explain the role of the learning rate in the perceptron algorithm. What happens if it is too high or too low?

- **Too high** → weights change drastically, algorithm may oscillate and fail to converge.
- **Too low** → learning is very slow, requiring many epochs to reach a good solution.
- Proper learning rate ensures faster convergence with stable results.

3. In real-world applications (e.g., spam detection, medical diagnosis), where could a perceptron be applied successfully?

Perceptrons can be applied in:

- **Spam detection** → classify emails as spam/ham.
- **Medical diagnosis** → classify whether a tumor is malignant/benign.
- **Sentiment analysis** → classify text as positive/negative.
- **Fraud detection** → classify transactions as fraudulent or normal.